

# PyMOR Usage

## Usage Examples: Modifying your Compute Resources

Paul Gierz

# Motivation

- Depending on the amount of data you have to process, it might be desirable to scale your computing resources
- From the user config, you can very easily:
  - Switch between processing backends (Dask vs. Dask-Jobqueue)
  - Scale how many SLURM worker jobs start
    - Other Jobqueue systems (e.g. PBS) are not yet supported!
  - Disable Dask entirely (good for debugging)

# User Config Settings

- `pymor.dask_cluster` can be `slurm` or `local`
- `pymor.dask_cluster_scaling_mode` can be `fixed` or `adapt`
- In adaptive mode, Dask (and SLURM) will automatically add or remove workers for you based upon how much number-crunching you need to do

```
1  general:
2      cmor_version: CMIP6
3      CMIP_Tables_Dir: /work/ab0995/a270243/pymor_workshop/cmip6-cmor-tables/Tables/
4      CV_Dir: /work/ab0995/a270243/pymor_workshop/cmip6-cmor-tables/CMIP6_CVs/
5  pymor:
6      warn_on_no_rule: False
7      dask_cluster: slurm
8      dask_cluster_scaling_mode: fixed # TODO: Change this during the exercise
9      dask_cluster_scaling_fixed_jobs: 1 #TODO: Change this during the exercise
10
11  rules:
12      - name: "linear trend example"
13          cmor_variable: tas
14          experiment_id: "piControl"
```

# User Config Settings

You can further specify settings for Dask in your configuration yaml!

See Dask Documentation for more:

- <https://docs.dask.org/en/stable/configuration.html>
- <https://docs.dask.org/en/stable/configuration.html#distributed-worker>
- <https://jobqueue.dask.org/en/latest/clusters-configuration.html>

```

25
26 # Settings for using dask-distributed
27 distributed:
28     worker:
29         memory:
30             target: 0.6 # Target 60% of worker memory usage
31             spill: 0.7 # Spill to disk when 70% of memory is used
32             pause: 0.8 # Pause workers if memory usage exceeds 80%
33             terminate: 0.95 # Terminate workers at 95% memory usage
34         resources:
35             CPU: 4 # Assign 4 CPUs per worker
36             death-timeout: 600 # Worker timeout if no heartbeat (seconds): Keep workers alive for 5
37 # SLURM-specific settings for launching workers
38 jobqueue:
39     slurm:
40         name: pymorize-worker
41         queue: compute # SLURM queue/partition to submit jobs
42         account: ab0995 # SLURM project/account name
43         cores: 4 # Number of cores per worker
44         memory: 128GB # Memory per worker
45         walltime: '00:30:00' # Maximum walltime per job
46         interface: ib0 # Network interface for communication
47         job-extra-directives: # Additional SLURM job options
48             - '--exclusive' # Run on exclusive nodes
49             - '--nodes=1'
50         # Worker template
51         worker-extra:
52             - "--nthreads"
53             - 4
54             - "--memory-limit"
55             - "128GB"
56             - "--lifetime"
57             - "25m"
58             - "--lifetime-stagger"
59             - "4m"
60         # How to launch workers and scheduler
61         job-cpu: 128
62         job-mem: 256GB
63         # worker-command: dask-worker
64         processes: 32 # Limited by memory per worker!
65         # scheduler-command: dask-scheduler

```

# Exercises

1. Run one of your examples locally if you have been using SLURM. What changes?
2. Try adding more worker nodes by increasing the `dask_cluster_scaling_fixed_jobs` field